

Project "Go Live" Technical Report

Executive Summary

This report details the technical architecture, deployment strategy, and critical steps required to take the **AutoGenCourseProject** (Course Generator) significantly from a local development application to a live, production-ready system on Google Cloud Platform (GCP).

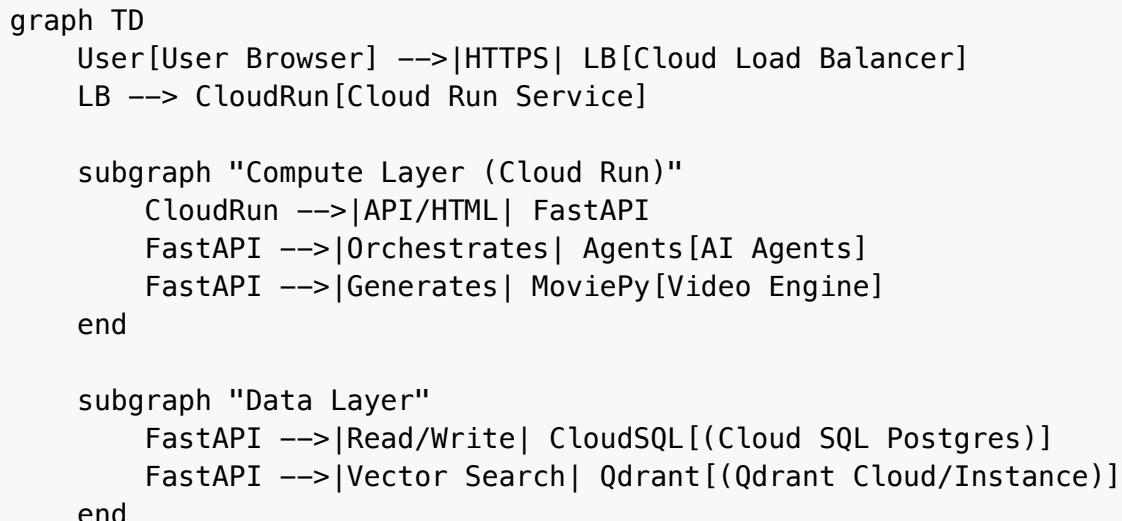
The system currently leverages a FastAPI backend, AI agents for content generation, and a local PostgreSQL database. The transition to production will involve containerization, managed cloud services for scalability, and a shift to distributed storage for generated media.

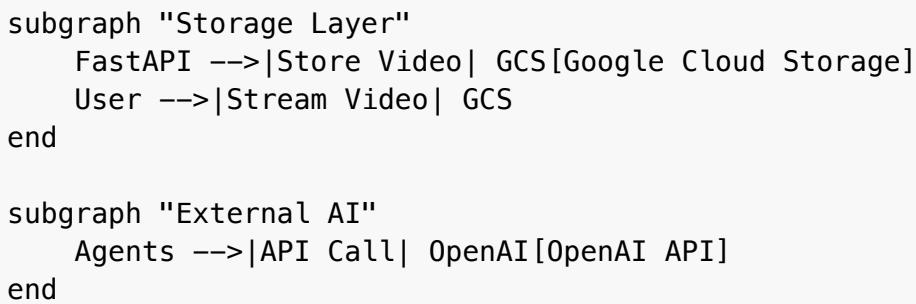
1. System Architecture Overview

1.1 Technology Stack

- **Backend Framework:** FastAPI (Python 3.10+)
 - *Role:* HTTP API, Request Handling, AI Orchestration.
- **Database:** PostgreSQL 15+ (Production: Cloud SQL)
 - *Role:* Structured storage for Users, Courses, Modules, Lessons, and Enrollments.
- **AI & Logic:**
 - **LLM Interface:** OpenAI API (`gpt-4o-mini`).
 - **Agent Framework:** AutoGen (AgentManager) for multi-agent workflows.
 - **Vector Search:** Qdrant (RAG memory & context retrieval).
- **Media Generation:**
 - **Video:** MoviePy (Programmatic video editing).
 - **Images:** Pillow (Image manipulation).
- **Frontend:**
 - Server-Side Rendering (SSR) with Jinja2 templates.
 - Vanilla CSS + JavaScript for interactivity (Course Builder, Dashboard).

1.2 Architecture Diagram





2. Deployment Strategy

The production environment will utilize **Google Cloud Platform (GCP)** largely for its strong support for containerized serverless applications and managed database services.

2.1 Core Infrastructure

- **Compute: Google Cloud Run.**
 - *Why:* Serverless, scales to zero (costs effectively zero when not in use), handles concurrent requests efficiently, and simplifies container deployment.
- **Database: Cloud SQL for PostgreSQL.**
 - *Why:* Fully managed, automated backups, high availability.
- **Storage: Google Cloud Storage (GCS) buckets.**
 - *Why:* Cloud Run has an ephemeral filesystem (files are lost on restart). Generated videos MUST be stored in GCS to be persistent and accessible via public URLs.

2.2 Continuous Integration/Deployment (CI/CD)

- **Tool: Google Cloud Build.**
- **Trigger:** Push to `main` branch.
- **Process:**
 1. Build Docker image.
 2. Push image to Google Artifact Registry.
 3. Deploy new revision to Cloud Run.
 4. Run database migrations (`alembic upgrade head`).

3. Critical "Go Live" Technical Changes

Before the application can be successfully deployed, the following technical refactors are **mandatory**:

3.1 Migration to Cloud Storage (Major Priority)

- **Current State:** Videos are saved to `course_material_service/static/videos` (local disk).
- **Issue:** On Cloud Run, this directory is temporary. A user might generate a video, but when they try to view it 5 minutes later, it could be gone if the instance reset.
- **Solution:** Refactor `course_material_service/video_builder.py` and `routes/generate_course.py`.
 - Implement a `StorageService` interface.

- After **MoviePy** generates the file to `/tmp`, upload it immediately to a GCS bucket (e.g., `gs://autogen-course-assets/videos/`).
- Save the **Public URL** (Signed URL or Public Bucket URL) in the database instead of the local filename.

3.2 Asynchronous Task Handling

- **Current State:** Video generation happens inside the HTTP request (using `run_in_threadpool`).
- **Issue:** Video rendering is CPU-heavy and slow. Cloud Run has a request timeout (default 5 mins). If a video takes 6 minutes, the request fails and the process is killed.
- **Solution:**
 - **Phase 1 (MVP):** Increase Cloud Run timeout to 15-30 minutes (doable via flags).
 - **Phase 2 (Robust):** Use **Cloud Tasks** or **Pub/Sub**.
 1. User clicks "Generate".
 2. App sends message to Queue.
 3. A separate "Worker" instance (or the same API with a background processor) picks up the job.
 4. When finished, it updates the DB status to "Completed".
 5. Frontend polls for status or uses WebSockets.

3.3 Database Connection Pooling

- **Requirement:** Ensure **SQLAlchemy** is configured with `pool_size` and `max_overflow` appropriate for Cloud Run concurrency to avoid exhausting Cloud SQL connections.

4. Operational "How It's Gonna Work"

4.1 The User Experience

1. **Login:** User logs in (Auth state stored in encrypted HTTP-only cookies).
2. **Course Design:** User interacts with the Course Builder. All changes are saved instantly to the Cloud SQL database.
3. **Generation:** User clicks "Generate Video".
 - The browser shows a "Processing..." progress bar.
 - The backend orchestrates the AI agents to write the script.
 - The backend renders the video and uploads it to Cloud Storage.
4. **Playback:** The video player in the dashboard streams the content directly from the global CDN (Cloud Storage) for low latency.

4.2 Monitoring & Maintenance

- **Logs:** All application logs (FastAPI info, Errors) are automatically captured by **Cloud Logging**.
- **Performance:** **Cloud Trace** will show latency bottlenecks (e.g., if OpenAI API is slow).
- **Cost Management:** Set up Budget Alerts in GCP billing to prevent unexpected costs from high GPU/CPU usage during video rendering.

5. Cost Estimates (Approximate)

WE ARE USING FREE LIMITS OF GOOGLE CLOUD

Service	Estimated Cost (Low Usage)	Notes
Cloud Run	\$0 - \$5 / month	Free tier covers 2M requests. Pays for CPU/Time.
Cloud SQL	~\$10 - \$30 / month	Micro instance. The main fixed cost.
Cloud Storage	< \$1 / month	Cheap for storing mp4 files.
OpenAI API	Usage Based	Depends on <code>gpt-4o-mini</code> token count (~\$0.15/million input tokens).
Qdrant Cloud	Free Tier / \$25+	Free tier available for small collections.

6. Next Steps Checklist

- **Step 1:** Create GCP Project and enable APIs (Run, SQL, Build, Storage).
- **Step 2:** Create a GCS Bucket for assets.
- **Step 3:** Refactor code to upload generated videos to GCS.
- **Step 4:** Provision Cloud SQL PostgreSQL instance and create DB.
- **Step 5:** Set up `cloudbuild.yaml` for automated deployment.
- **Step 6:** Deploy and validate end-to-end flow.